

DFS:=

```
#include <iostream>

#include <vector>

#include <stack>

#include <omp.h>

using namespace std;

const int MAX = 100000;

vector<int> graph[MAX];

bool visited[MAX];

void dfs(int node) {

    stack<int> s;

    s.push(node);

    while (!s.empty()) {

        int curr_node = s.top();

        if (!visited[curr_node]) {

            visited[curr_node] = true;

            s.pop();

            cout<<curr_node<<" ";

            #pragma omp parallel for

            for (int i = 0; i < graph[curr_node].size(); i++) {

                int adj_node = graph[curr_node][i];

                if (!visited[adj_node]) {

                    s.push(adj_node);

                }

            }

        }

    }

}

int main() {

    int n, m, start_node;

    cout<<"Enter no. of Node,no. of Edges and Starting Node of graph:\n";

    cin >> n >> m >> start_node;

    //n: node,m:edges

    cout<<"Enter pair of node and edges:\n";

    for (int i = 0; i < m; i++) {
```

```
int u, v;

cin >> u >> v;

//u and v: Pair of edges
graph[u].push_back(v);
graph[v].push_back(u);
}

#pragma omp parallel for
for (int i = 0; i < n; i++) {
    visited[i] = false;
}

dfs(start_node);

return 0;
}
```

BFS :=

```
#include<iostream>

#include<stdlib.h>

#include<queue>

using namespace std;

class node

{

public:

node *left, *right;

int data;

};

class Breadthfs

{

public:

node *insert(node *, int);

void bfs(node *);

};

node *insert(node *root, int data)

// inserts a node in tree

{

if(!root)

{

root=new node;

root->left=NULL;

root->right=NULL;

root->data=data;

return root;

}

queue<node *> q;

q.push(root);

while(!q.empty())

{

node *temp=q.front();

q.pop();

if(temp->left==NULL)
```

```

{
temp->left=new node;
temp->left->left=NULL;
temp->left->right=NULL;
temp->left->data=data;
return root;
}
else
{
q.push(temp->left);
}
if(temp->right==NULL)
{
temp->right=new node;
temp->right->left=NULL;
temp->right->right=NULL;
temp->right->data=data;
return root;
}
else
{
q.push(temp->right);
}
}
}

void bfs(node *head)
{
queue<node*> q;
q.push(head);
int qSize;
while (!q.empty())
{
qSize = q.size();
#pragma omp parallel for
//creates parallel threads

```

```

for (int i = 0; i < qSize; i++)
{
    node* currNode;
    #pragma omp critical
    {
        currNode = q.front();
        q.pop();
        cout<<"\t"<<currNode->data;
        }// prints parent node
    #pragma omp critical
    {
        if(currNode->left)// push parent's left node in queue
            q.push(currNode->left);
        if(currNode->right)
            q.push(currNode->right);
        }// push parent's right node in queue
    }
}

int main(){
    node *root=NULL;

    int data;

    char ans;

    do
    {
        cout<<"\n enter data=>";

        cin>>data;

        root=insert(root,data);

        cout<<"do you want insert one more node?";

        cin>>ans;

    }while(ans=='y' || ans=='Y');

    bfs(root);

    return 0;
}

```

Paralle Bubble Sort :=

```
import time
import random

start = time.perf_counter()

def Parallel_Bubble_sort(lst):
    Sorted = 0
    n = len(lst)
    while Sorted == 0:
        Sorted = 1
        for i in range(0, n-1, 2):
            if lst[i] > lst[i+1]:
                lst[i], lst[i+1] = lst[i+1], lst[i]
                Sorted = 0
        for i in range(1, n-1, 2):
            if lst[i] > lst[i+1]:
                lst[i], lst[i+1] = lst[i+1], lst[i]
                Sorted = 0
    print(lst)

lst = [(random.randint(0,100)) for i in range(100)]
Parallel_Bubble_sort(lst)
finish = time.perf_counter()
print(f'Finished in {round(finish-start,2)} second(s)')
```

MERGE SORT :=

```
def merge(arr, l, m, r):  
  
    n1 = m - l + 1  
  
    n2 = r - m  
  
    # create temp arrays  
  
    L = [0] * (n1)  
  
    R = [0] * (n2)  
  
    for i in range(0, n1):  
  
        L[i] = arr[l + i]  
  
    for j in range(0, n2):  
  
        R[j] = arr[m + 1 + j]  
  
    i = 0 # Initial index of first subarray  
  
    j = 0 # Initial index of second subarray  
  
    k = l # Initial index of merged subarray  
  
    while i < n1 and j < n2:  
  
        if L[i] <= R[j]:  
  
            arr[k] = L[i]  
  
            i += 1  
  
        else:  
  
            arr[k] = R[j]  
  
            j += 1  
  
            k += 1  
  
  
    while i < n1:  
  
        arr[k] = L[i]  
  
        i += 1  
  
        k += 1  
  
  
    while j < n2:  
  
        arr[k] = R[j]  
  
        j += 1  
  
        k += 1  
  
  
def mergeSort(arr, l, r):  
  
    if l < r:
```

```
m = l+(r-l)//2
mergeSort(arr, l, m)
mergeSort(arr, m+1, r)
merge(arr, l, m, r)
```

```
# Driver code to test above
```

```
arr = [12, 11, 13, 5, 6, 7]
```

```
n = len(arr)
```

```
print("Given array is")
```

```
for i in range(n):
```

```
    print("%d" % arr[i],end=" ")
```

```
mergeSort(arr, 0, n-1)
```

```
print("\n\nSorted array is")
```

```
for i in range(n):
```

```
    print("%d" % arr[i],end=" ")
```


Implement Min, Max, Sum and Average operations using Parallel Reduction :=

```
#include <iostream>

#include <vector>

#include <omp.h>

#include <climits>

using namespace std;

void min_reduction(vector<int>& arr) {

    int min_value = INT_MAX;

    #pragma omp parallel for reduction(min: min_value)

    for (int i = 0; i < arr.size(); i++) {

        if (arr[i] < min_value) {

            min_value = arr[i];

        }

    }

    cout << "Minimum value: " << min_value << endl;

}

void max_reduction(vector<int>& arr) {

    int max_value = INT_MIN;

    #pragma omp parallel for reduction(max: max_value)

    for (int i = 0; i < arr.size(); i++) {

        if (arr[i] > max_value) {

            max_value = arr[i];

        }

    }

    cout << "Maximum value: " << max_value << endl;

}

void sum_reduction(vector<int>& arr) {

    int sum = 0;

    #pragma omp parallel for reduction(+: sum)

    for (int i = 0; i < arr.size(); i++) {

        sum += arr[i];

    }

    cout << "Sum: " << sum << endl;

}
```

```
void average_reduction(vector<int>& arr) {  
    int sum = 0;  
    #pragma omp parallel for reduction(+: sum)  
    for (int i = 0; i < arr.size(); i++) {  
        sum += arr[i];  
    }  
    cout << "Average: " << (double)sum / arr.size() << endl;  
}  
  
int main() {  
    vector<int> arr;  
    arr.push_back(5);  
    arr.push_back(2);  
    arr.push_back(9);  
    arr.push_back(1);  
    arr.push_back(7);  
    arr.push_back(6);  
    arr.push_back(8);  
    arr.push_back(3);  
    arr.push_back(4);  
    min_reduction(arr);  
    max_reduction(arr);  
    sum_reduction(arr);  
    average_reduction(arr);  
}
```

Linear Regression :=

```
import pandas as pd

# Load the dataset from a CSV file
df = pd.read_csv('/content/HousingData.csv')

# Display the first few rows of the dataset
print(df.head())

#Step 2: Preprocess the data
from sklearn.preprocessing import StandardScaler # Split the data into input and output variables
X = df.drop('MEDV', axis=1)
y = df['MEDV']

# Scale the input features
scaler = StandardScaler()
X = scaler.fit_transform(X)

# Display the first few rows of the scaled input features print(X[:5])

#Step 3: Split the dataset
from sklearn.model_selection import train_test_split # Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Print the shapes of the training and testing sets
print('Training set shape:', X_train.shape, y_train.shape)
print('Testing set shape:', X_test.shape, y_test.shape)

#Step 4: Define the model architecture
from keras.models import Sequential
from keras.layers import Dense, Dropout

# Define the model architecture
model = Sequential()
model.add(Dense(64, input_dim=13, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(32, activation='relu'))
model.add(Dense(1))

# Display the model summary
print(model.summary())

#Step 5: Compile the model

# Compile the model
model.compile(loss='mean_squared_error', optimizer='adam', metrics=['mean_absolute_error'])
```

#Step 6: Train the model

```
from keras.callbacks import EarlyStopping # Train the model
```

```
early_stopping = EarlyStopping(monitor='val_loss', patience=5)
```

```
history = model.fit(X_train, y_train, validation_split=0.2, epochs=100, batch_size=32, callbacks=[early_stopping])
```

Plot the training and validation loss over epochs

```
import matplotlib.pyplot as plt
```

```
plt.plot(history.history['loss'])
```

```
plt.plot(history.history['val_loss'])
```

```
plt.title('Model Loss')
```

```
plt.xlabel('Epochs')
```

```
plt.ylabel('Loss')
```

```
plt.legend(['Training', 'Validation'])
```

```
plt.show()
```

#Step 7: Evaluate the model

Evaluate the model on the testing set

```
loss, mae = model.evaluate(X_test, y_test)
```

Print the mean absolute error

```
print('Mean Absolute Error:', mae)
```

Movie Review Classification using IMDB Dataset :-

```
import numpy as np

from keras.datasets import imdb

from keras.preprocessing.sequence import pad_sequences

from keras.models import Sequential

from keras.layers import Embedding, Bidirectional, LSTM, Dense # Load the IMDB dataset

(x_train, y_train), (x_test, y_test) = imdb.load_data()

# Pad or truncate the sequences to a fixed length of 250 words

max_len = 250

x_train = pad_sequences(x_train, maxlen=max_len)

x_test = pad_sequences(x_test, maxlen=max_len)

# Define the deep neural network architecture

model = Sequential()

model.add(Embedding(input_dim=10000, output_dim=128, input_length=max_len))

model.add(Bidirectional(LSTM(64, return_sequences=True)))

model.add(Bidirectional(LSTM(32)))

model.add(Dense(1, activation='sigmoid')) # Compile the model

model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy']) # Train the model

history = model.fit(x_train, y_train, epochs=10, batch_size=128, validation_split=0.2) # Evaluate the model on the test set

loss, acc = model.evaluate(x_test, y_test, batch_size=128)

print(f'Test accuracy: {acc:.4f}, Test loss: {loss:.4f}')
```

CNN MNIST Fashhion Dataset :=

```
import tensorflow as tf

from tensorflow import keras

import numpy as np

import matplotlib.pyplot as plt # Load the dataset

fashion_mnist = keras.datasets.fashion_mnist

(train_images, train_labels), (test_images, test_labels) = fashion_mnist.load_data() # Normalize the images

train_images = train_images / 255.0

test_images = test_images / 255.0

# Define the model

model = keras.Sequential([ keras.layers.Flatten(input_shape=(28, 28)), keras.layers.Dense(128, activation='relu'),
keras.layers.Dense(10, activation='softmax')])

# Compile the model

model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

# Train the model

model.fit(train_images, train_labels, epochs=10) # Evaluate the model

test_loss, test_acc = model.evaluate(test_images, test_labels)

print('Test accuracy:', test_acc)

# Make predictions

predictions = model.predict(test_images)

predicted_labels = np.argmax(predictions, axis=1)

# Show some example images and their predicted labels

num_rows = 5

num_cols = 5

num_images = num_rows * num_cols

plt.figure(figsize=(2 * 2 * num_cols, 2 * num_rows))

for i in range(num_images):

    plt.subplot(num_rows, 2 * num_cols, 2 * i + 1)

    plt.imshow(test_images[i], cmap='gray')

    plt.axis('off')

    plt.subplot(num_rows, 2 * num_cols, 2 * i + 2)

    plt.bar(range(10), predictions[i])

    plt.xticks(range(10))

    plt.ylim([0, 1])
```

```
plt.tight_layout()
plt.title(f"Predicted label: {predicted_labels[i]}")
plt.show()
```

CUDA Program for Addition of Two Large Vectors:

```
#include <stdio.h> #include <stdlib.h>

// CUDA kernel for vector addition
global void vectorAdd(int *a, int *b, int *c, int n) { int i = blockIdx.x * blockDim.x + threadIdx.x; if (i <
n) {
c[i] = a[i] + b[i];
}
}

int main() {
int n = 1000000; // Vector size int *a, *b, *c; // Host vectors
int *d_a, *d_b, *d_c; // Device vectors int size = n * sizeof(int); // Size in bytes
// Allocate memory for host vectors a = (int*) malloc(size);
b = (int*) malloc(size); c = (int*) malloc(size);
// Initialize host vectors for (int i = 0; i < n; i++) {
a[i] = i;
b[i] = i;}
// Allocate memory for device vectors cudaMalloc((void**) &d_a, size);
cudaMalloc((void**) &d_b, size); cudaMalloc((void**) &d_c, size);
// Copy host vectors to device vectors
cudaMemcpy(d_a, a, size, cudaMemcpyHostToDevice); cudaMemcpy(d_b, b, size,
cudaMemcpyHostToDevice);
// Define block size and grid size int blockSize = 256;
int gridSize = (n + blockSize - 1) / blockSize;
// Launch kernel
vectorAdd<<<gridSize, blockSize>>>(d_a, d_b, d_c, n);
// Copy device result vector to host result vector cudaMemcpy(c, d_c, size,
cudaMemcpyDeviceToHost);
// Verify the result
for (int i = 0; i < n; i++) { if (c[i] != 2*i) {
printf("Error: c[%d] = %d\n", i, c[i]); break;
}
}
// Free device memory cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);
// Free host memory free(a); free(b);
free(c);
```



```
return 0;
```

```
}
```

Cuda Program for matrix multiplication :=

```
#include <stdio.h> #define BLOCK_SIZE 16

global void matrix_multiply(float *a, float *b, float *c, int n)
{
    int row = blockIdx.y * blockDim.y + threadIdx.y; int col = blockIdx.x * blockDim.x + threadIdx.x; float
    sum = 0;

    if (row < n && col < n) { for (int i = 0; i < n; ++i) {
        sum += a[row * n + i] * b[i * n + col];
    }
    c[row * n + col] = sum;}
}

int main(){
    int n = 1024;

    size_t size = n * n * sizeof(float); float *a, *b, *c;
    float *d_a, *d_b, *d_c; cudaEvent_t start, stop; float elapsed_time;

    // Allocate host memory
    a = (float*)malloc(size);
    b = (float*)malloc(size);
    c = (float*)malloc(size);

    // Initialize matrices
    for (int i = 0; i < n * n; ++i) { a[i] = i % n;
        b[i] = i % n;
    }

    // Allocate device memory
    cudaMalloc(&d_a, size);
    cudaMalloc(&d_b, size);
    cudaMalloc(&d_c, size);

    // Copy input data to device
    cudaMemcpy(d_a, a, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_b, b, size, cudaMemcpyHostToDevice);

    // Set kernel launch configuration
    dim3 threads(BLOCK_SIZE, BLOCK_SIZE);

    dim3 blocks((n + threads.x - 1) / threads.x, (n + threads.y - 1) / threads.y);

    // Launch kernel cudaEventCreate(&start); cudaEventCreate(&stop); cudaEventRecord(start);
    matrix_multiply<<<blocks, threads>>>(&d_a, &d_b, &d_c, n); cudaEventRecord(stop);
```

```
cudaEventSynchronize(stop); cudaEventElapsedTime(&elapsed_time, start, stop);  
  
// Copy output data to host  
  
cudaMemcpy(c, d_c, size, cudaMemcpyDeviceToHost);  
  
// Print elapsed time  
  
printf("Elapsed time: %f ms\n", elapsed_time);  
  
// Free device memory cudaFree(d_a); cudaFree(d_b); cudaFree(d_c);  
  
// Free host memory free(a); free(b);  
  
free(c);  
  
return 0;  
  
}
```